

Solution:

All below model is trained on 'wnut17train.conll', and tested on 'emerging.test.conll'

a) CRF trained with default setting obtained avg f1_score on test data of 0.916.

	precision	recall	f1-score	support
O	0.945	0.996	0.970	21934
B-corporation	0.000	0.000	0.000	117
I-corporation	0.000	0.000	0.000	24
B-creative-work	0.500	0.013	0.025	231
I-creative-work	0.786	0.046	0.088	237
B-group	0.000	0.000	0.000	97
I-group	0.143	0.026	0.043	39
B-location	0.235	0.189	0.209	122
I-location	0.179	0.128	0.149	39
B-person	0.516	0.093	0.158	355
I-person	0.515	0.168	0.254	101
B-product	0.667	0.036	0.068	56
I-product	0.250	0.024	0.043	42
accuracy			0.938	23394
macro avg	0.364	0.132	0.154	23394
weighted avg	0.913	0.938	0.916	23394

b) Softmax classifier trained with default setting, obtained avg f1_score on test data of 0.907, where CRF performed better with a higher f1_score. In part c, I will proceed to fine tune hyperparameters of the classifier.

	precision	recall	f1-score	support
O	0.938	1.000	0.968	21934
B-corporation	0.000	0.000	0.000	117
I-corporation	0.000	0.000	0.000	24
B-creative-work	0.000	0.000	0.000	231
I-creative-work	0.000	0.000	0.000	237
B-group	0.000	0.000	0.000	97
I-group	0.000	0.000	0.000	39
B-location	0.000	0.000	0.000	122
I-location	0.000	0.000	0.000	39
B-person	0.000	0.000	0.000	355
I-person	0.000	0.000	0.000	101
B-product	0.000	0.000	0.000	56
I-product	0.000	0.000	0.000	42
accuracy			0.938	23394
macro avg	0.072	0.077	0.074	23394
weighted avg	0.879	0.938	0.907	23394

c) **Softmax classifier hyperparameter tuning: test avg f1_score = 0.952**

I tried below values for each hyperparameter, and the avg f1 score on the dev dataset will be the criteria to choose the best parameter settings.


Table of hyperparameter tuning validation f1 score

Training config (parameters same as default other than indicated)	Avg f1 score on dev dataset (‘emerging.dev.conll’)
Default 'batch_size': 4, 'half_window': 2, 'embed_dim': 25, 'hidden_dim': 25, 'num_classes': 13, 'freeze_embeddings': False, 'lr': 0.02, 'epochs': 100 'Num_layers': 1	0.8824671395802548
Batch_size 2	0.8786100406439366
Batch_size 8	0.8824671395802548
Embed_dim 50	0.8824181390344367
Embed_dim 100 (2nd best)	0.8828044283450783
Embed_dim 200	0.8824118855465907
Embed_dim 300	0.8822393988274964
Half window 3	0.8824036834877962
Half window 4	0.8824963458400458
Half window 5	0.8824354125841017
Hidden dim 50	0.8824671395802548
Hidden dim 100	0.8824631384713871
Number layer 2	0.8823476631930625
Number layer 3(best)	0.8828561475496695
Number layer 4	0.8823008878819346
Freeze_embeddings True	0.8823326493950854
Lr 0.002	0.8821357462243560
Lr 0.1	0.8814530901515277
Epochs 200	0.8829305762387519
Epochs 300(improve with epochs)	0.8831023763895041
Combine two parameters Embed dim 100, num layer 3	0.8820091892966707

From the above table, the most effective parameter is 3 hidden layers(default 1) plus more epochs. To get better results, I also used a learning rate scheduler to reduce the rate by half per 100 epochs. The model was trained for 1000 epochs(when loss converges) and got avg f1 score of 0.952 on test data.

The learning curve and classification report are attached below.

Training curves



	precision	recall	f1-score	support
O	0.968	0.999	0.983	44296
B-corporation	0.778	0.356	0.488	177
I-corporation	0.000	0.000	0.000	36
B-creative-work	0.471	0.075	0.129	107
I-creative-work	0.743	0.310	0.437	168
B-group	0.308	0.027	0.050	148
I-group	0.444	0.037	0.068	108
B-location	0.714	0.389	0.503	391
I-location	0.741	0.244	0.367	164
B-person	0.619	0.261	0.367	468
I-person	0.727	0.213	0.330	225
B-product	0.375	0.030	0.056	100
I-product	0.833	0.062	0.115	81
accuracy			0.963	46469
macro avg	0.594	0.231	0.299	46469
weighted avg	0.952	0.963	0.952	46469